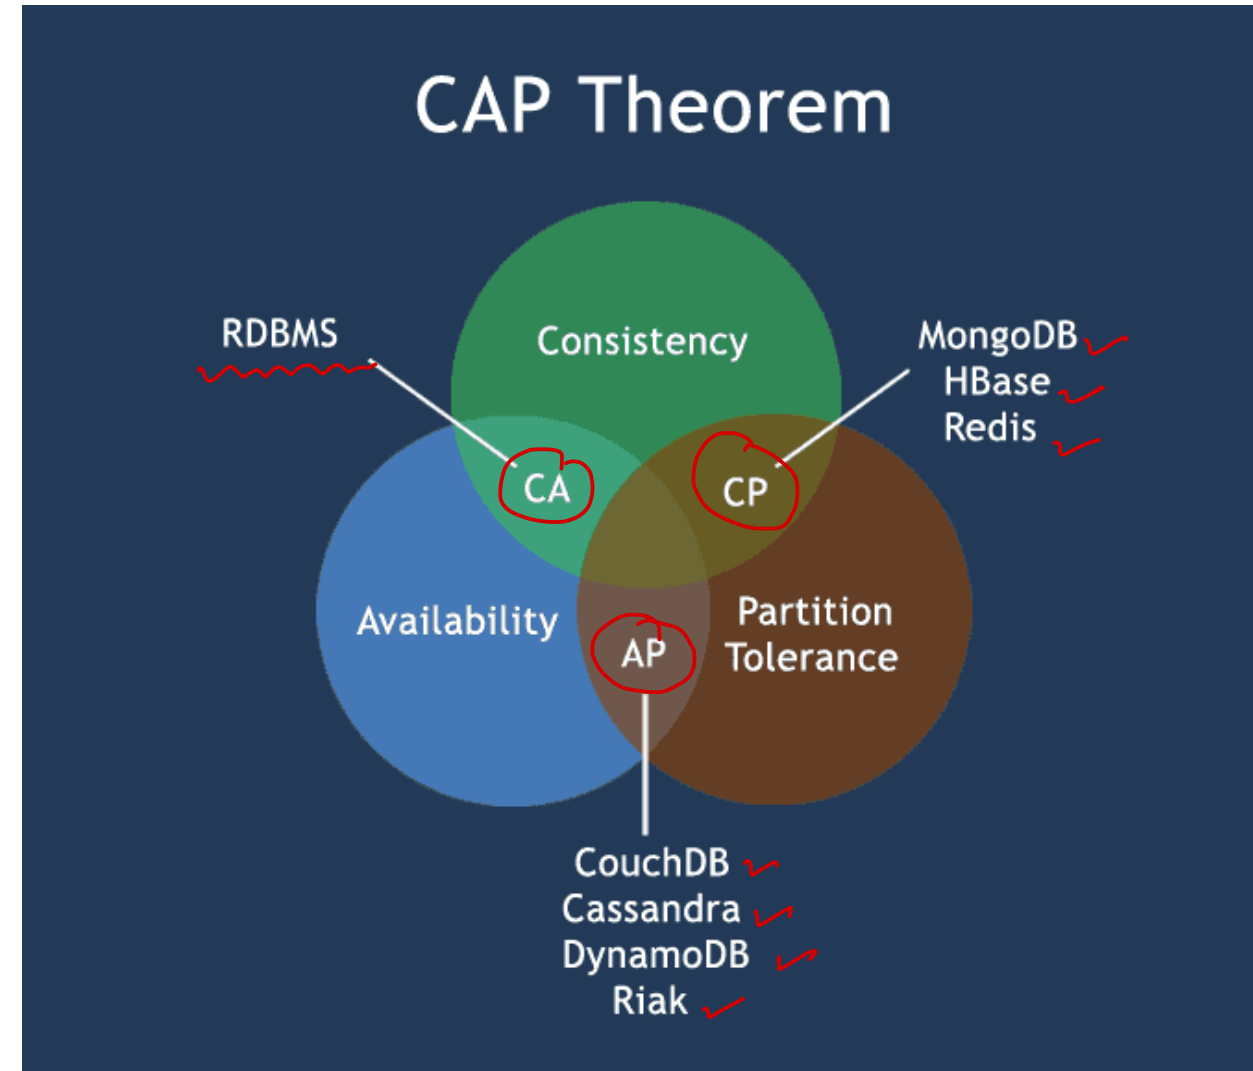



NoSQL Databases

Trainer: Mr. Nilesh Ghule

CAP (Brewer's) Theorem

- **Consistency** - Data is consistent after operation. After an update operation, all clients see the same data.
- **Availability** - System is always on (i.e. service guarantee), no downtime.
- **Partition Tolerance** - System continues to function even the communication among the servers is unreliable.
- **Brewer's Theorem**
 - It is impossible for a distributed data store to simultaneously provide more than two out of the above three guarantees.





Applications

- When to use NoSQL?

- ✓• Large amount of data (TBs)
- ✓• Many Read/Write ops
- ✓• Economical Scaling
- ✓• Flexible schema

- Examples:

- ✓• Social media
- ✓• Recordings
- ✓• Geospatial analysis
- ✓• Information processing

- When Not to use NoSQL?

- ✓• Need ACID transactions
- ✓• Fixed multiple relations
- ✓• Need joins
- ✓• Need high consistency

- Examples

- ✓• Financial transactions
- ✓• Business operations



RDBMS vs NoSQL

	RDBMS	NoSQL
Types	All types support <u>SQL standard</u>	Multiple types exists, such as <u>document stores</u> , <u>key value stores</u> , <u>column databases</u> , etc
History	Developed in <u>1970</u>	Developed in <u>2000s</u>
Examples	SQL Server, Oracle, MySQL ✓	MongoDB, HBase, Cassandra, Redis, Neo4J ✓
Data Storage Model	Data is stored in <u>rows</u> and <u>columns</u> in a <u>table</u> , where each column is of a specific type	The data model depends on the database type. It could be <u>Key-value pairs</u> , <u>documents</u> etc
Schemas	<u>Fixed structure and schema</u>	<u>Dynamic schema</u> . Structures can be accommodated
Scalability	<u>Scale up</u> approach is used (vertical)	<u>Scale out</u> approach is used (horizontal)
Transactions	Supports <u>ACID</u> and transactions	Supports partitioning and availability <u>BASE</u>
Consistency	<u>Strong consistency</u>	Dependent on the product [<u>Eventual Consistency</u>]
Support	High level of enterprise support ✓	Open source model ✓
Maturity	Have been around for a long time ✓	Some of them are mature; others are evolving ✓




NoSQL database

- NoSQL databases are non-relational. → ~~tabular~~
- There is no standardization/rules of how NoSQL database to be designed.
- All available NoSQL databases can be broadly categorized as follows:
 - ✓• Key-value databases
 - ✓• Column-oriented databases
 - ✓• Graph databases
 - ✓• Document oriented databases



Key-value database


- Based on Amazon's Dynamo database.
- For handling huge data of any type.
- Keys are unique and values can be of any type i.e. JSON, BLOB, etc.
- Implemented as big distributed hash-table for fast searching.
- Example: redis, dynamodb, riak, ...



Column-oriented databases


- Values of columns are stored contiguously.
- Better performance while accessing few columns and aggregations.
- Good for data-warehousing, business intelligence, CRM, ...
- Examples: hbase, cassandra, bigtable, ...

Select job, sum(sal) from emp
group by job;




RDBMS

columnar db



Graph databases

- Graph is collection of vertices and edges (lines connecting vertices).
- Vertices keep data, while edges represent relationships.
- Each node knows its adjacent nodes. Very good performance, when want to access all relations of an entity (irrespective of size of data).
- Examples: Neo4J, Titan, ...



Document oriented databases

- Document contains data as key-value pair as JSON or XML.
- Document schema is flexible & are added in collection for processing.
- RDBMS tables → Collections
- RDBMS rows → Documents
- RDBMS columns → Key-value pairs in document
- Examples: MongoDb, CouchDb, ...

Ex 1 JSON → Java Script Object Notation.

```
{  
  "id": 1,           → int  
  "title": "Let us C", → string  
  "author": "Karnetkar",  
  "price": 240.4     → double  
}
```

Ex 2 → JSON document

```
{  
  id: 1,  
  name: "Nilesh",  
  age: 38,  
  hobbies: ["Programming", "Reading", ...],  
  addr: { area: "Kotraj", city: "Pune", pin: 4110463},  
  political: false,  
  height: 5.9,  
  bloodgroup: null  
}
```

3





MongoDb Databases

Trainer: Mr. Nilesh Ghule



Agenda

- Introduction
- Installation
- JSON vs BSON
- Basic CRUD operations



Mongo Db

- Developed by 10gen in 2007
- Publicly available in 2009
- Open-source database which is controlled by 10gen
- Document oriented database → stores JSON documents
- Stores data in binary JSON. (BSON)
- Design Philosophy
 - MongoDB wasn't designed in a lab and is instead built from the experiences of building large scale, high availability, and robust systems.




JSON

- Java Script Object Notation
- Hierarchical way of organizing data
- Mongo stores JSON data into Binary form.



Mongo Server and Client

- MongoDB server (mongod) is developed in C, C++ and JS.
- MongoDB data is accessed via multiple client tools
 - ✓ mongo : client shell (JS). *obj. method (args) ;*
 - ✓ mongofiles : stores larger files in GridFS.
 - ✓ mongoimport / mongoexport : tools for data import / export.
 - ✓ mongodump / mongorestore : tools for backup / restore.
- MongoDb data can be accessed in application through client drivers available for all major programming languages e.g. Java, Python, Ruby, PHP, Perl, ...
- Mongo shell is follows JS syntax and allow to execute JS scripts.



MongoDb: Data Types

data	bson	values
null	10	
boolean	8	true, false
number	1 / 16 / 18	123, 456.78, NumberInt("24"), NumberLong("28")
string	2	"...."
date	9	new Date(), ISODate("yyyy-mm-ddThh:mm:ss")
array	4	[..., ..., ..., ...]
object	3	{ ... }



Mongo - INSERT

- show databases;
- use database;
- db.contacts.insert({name: "nilesh", mobile: "9527331338"});
- db.contacts.insertMany([
 {name: "nilesh", mobile: "9527331338"},
 {name: "nitin", mobile: "9881208115"}
]);
- Maximum document size is 16 MB.
- For each object unique id is generated by client (if _id not provided).
 - 12 byte unique id :: [counter(3) | pid(2) | machine(3) | timestamp(4)]
 - client process id
 - client
 - Field/key



Mongo – QUERY

- `db.contacts.find();` → returns cursor on which following ops allowed:
 - `hasNext()`, `next()`, `skip(n)`, `limit(n)`, `count()`, `toArray()`, `forEach(fn)`, `pretty()`
- Shell restrict to fetch 20 records at once. Press "it" for more records.
- `db.contacts.find({ name: "nilesh" });`
- `db.contacts.find({ name: "nilesh" }, { _id:0, name:1 });`
- Relational operators: `$eq`, `$ne`, `$gt`, `$lt`, `$gte`, `$lte`, `$in`, `$nin`
- Logical operators: `$and`, `$or`, `$nor`, `$not`
- Element operators: `$exists`, `$type`
- Evaluation operators: `$regex`, `$where`, `$mod`
- Array operators: `$size`, `$elemMatch`, `$all`, `$slice`



Mongo – DELETE

- `db.contacts.remove(criteria);`
- `db.contacts.deleteOne(criteria);`
- `db.contacts.deleteMany(criteria);`
- `db.contacts.deleteMany({});` → delete all docs, but not collection
- `db.contacts.drop();` → delete all docs & collection as well : efficient



Mongo – UPDATE

- `db.contacts.update(criteria, newObj);`
- Update operators: `$set`, `$inc`, `$dec`, `$push`, `$each`, `$slice`, `$pull`
- In place updates are faster (e.g. `$inc`, `$dec`, ...) than setting new object. If new object size mismatch with older object, data files are fragmented.
- Update operators: `$addToSet`
- example: `db.contacts.update({ name: "peter" },`
- `{ $push : { mobile: { $each : ["111", "222"], $slice : -3 } } });`
- `db.contacts.update({ name: "t" }, { $set : { "phone" : "123" } }, true);`
 - If doc with given criteria is absent, new one is created before update.





Thank you!

Nilesh Ghule <nilesh@sunbeaminfo.com>

